
Plant Pathology Challenge

類神經網路導論期末報告



問題定義

傳統上使用人工的方式，消耗人力及費時。再者，對於疾病的錯誤判斷可能導致抗藥性的病原體產生，使得疾病難以被根治。此為CVPR2020的FGVC7研討會挑戰。





輸入資料格式

輸入

特徵(feature):

1822張 RGB 照片

答案

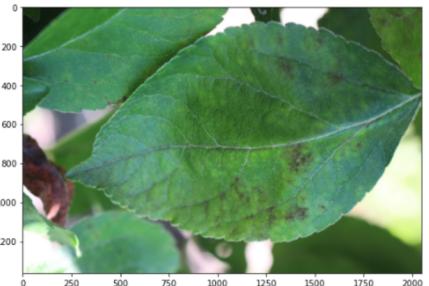
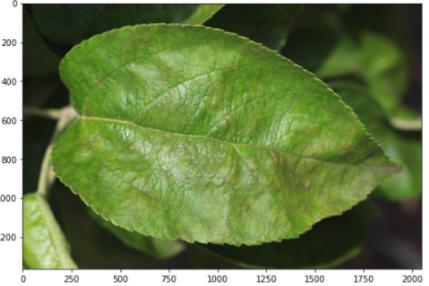
healthy、multiple_diseases、rust、scab

輸出

給定test資料集，所檢測的結果。

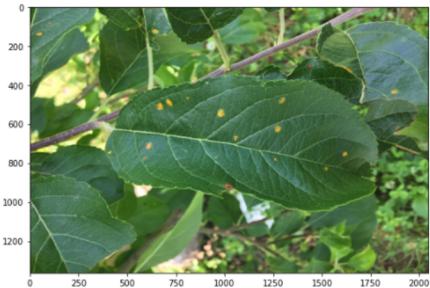
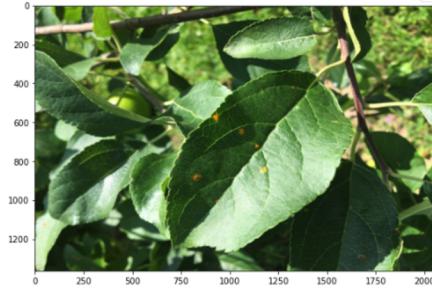
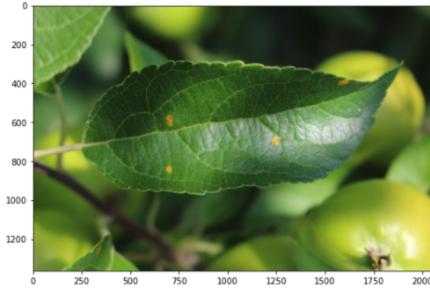
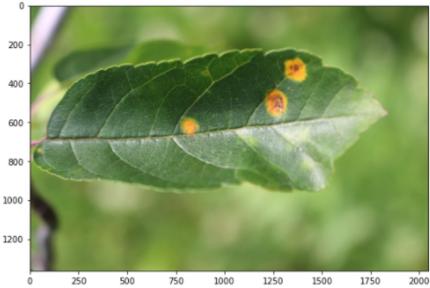


輸入答案 - scab

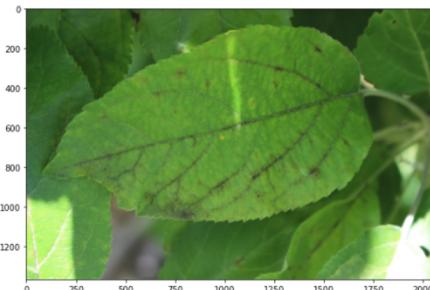
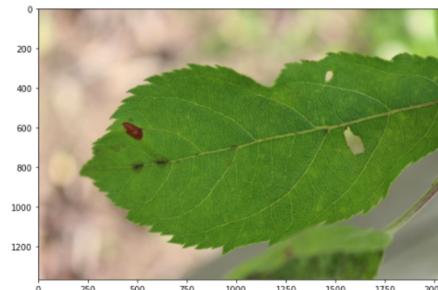
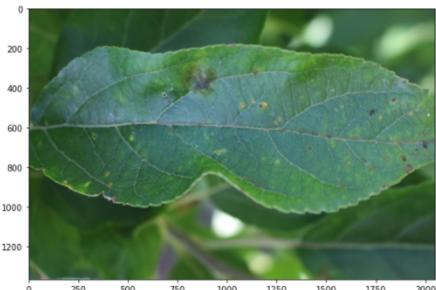
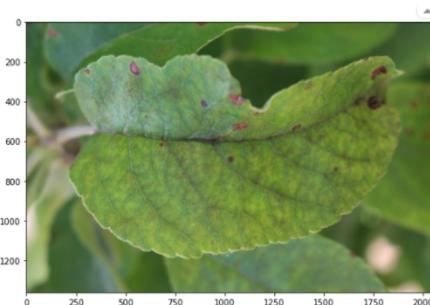
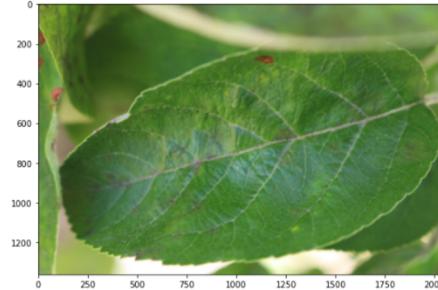
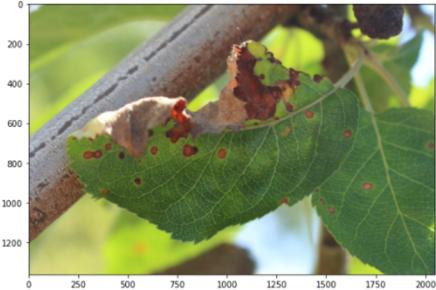




輸入答案 - rust

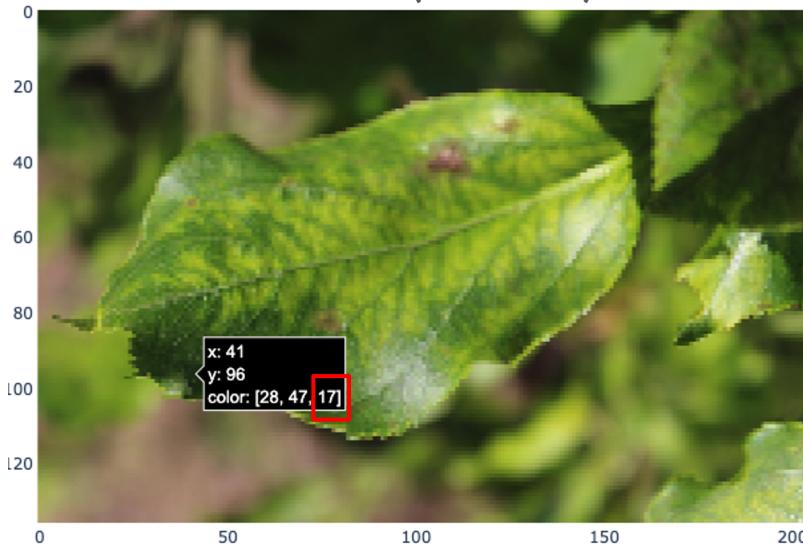


輸入答案 - multiple diseases

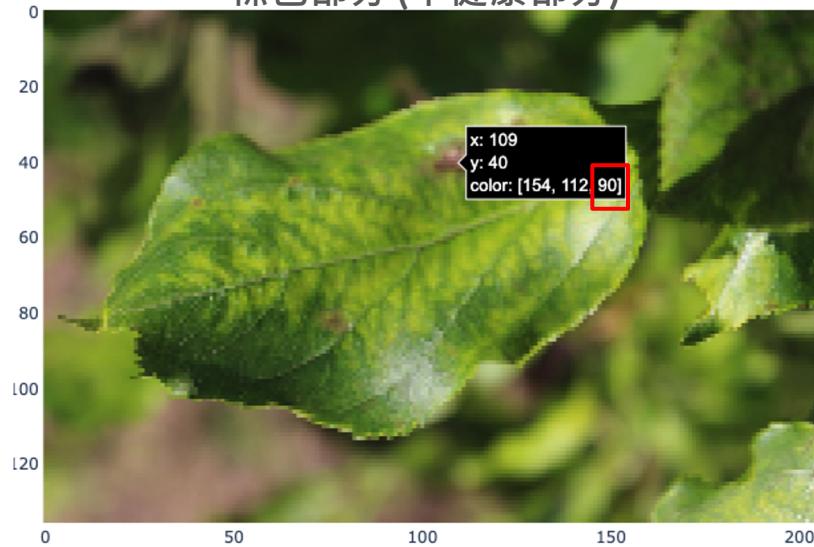


輸入資料視覺化

綠色部分 (健康部分)



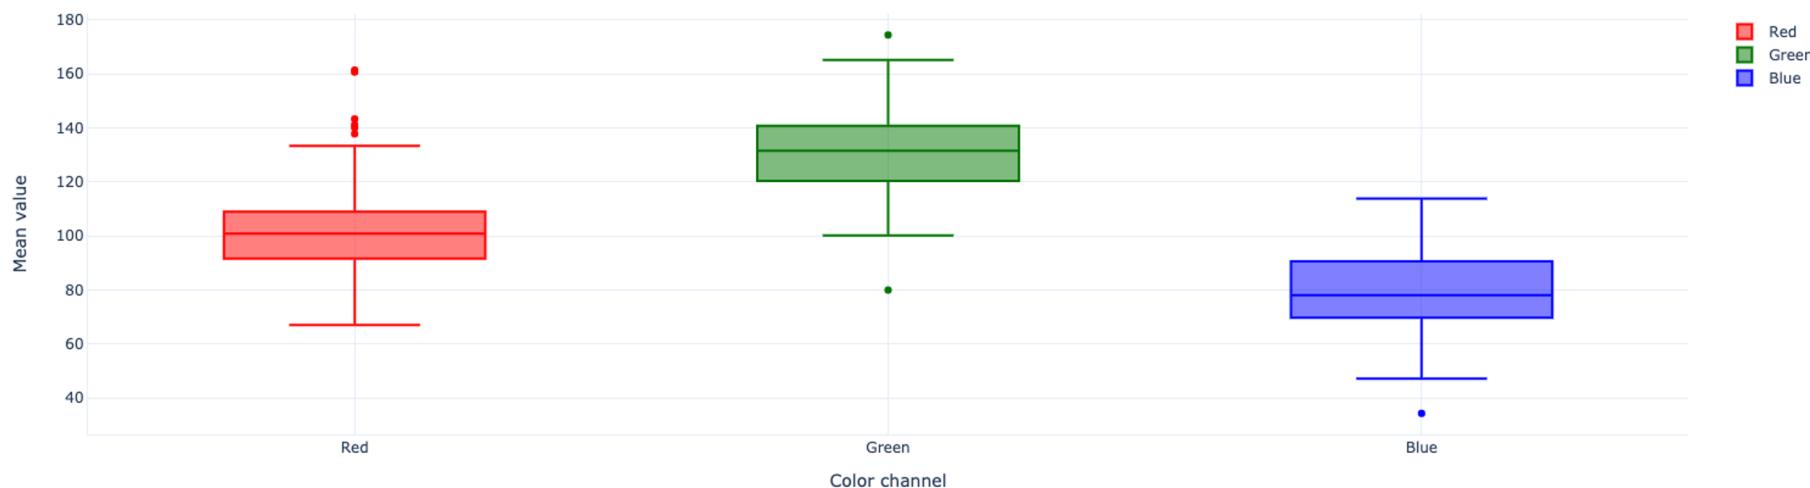
棕色部分 (不健康部分)





輸入資料統計

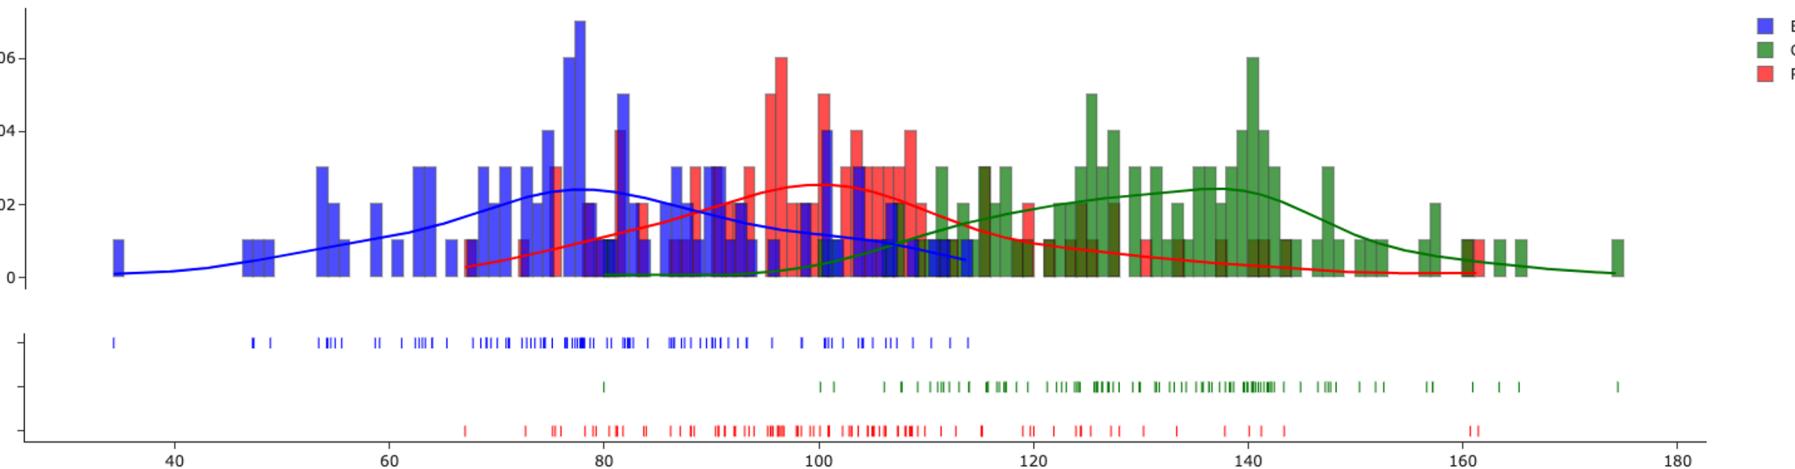
Mean value vs. Color channel





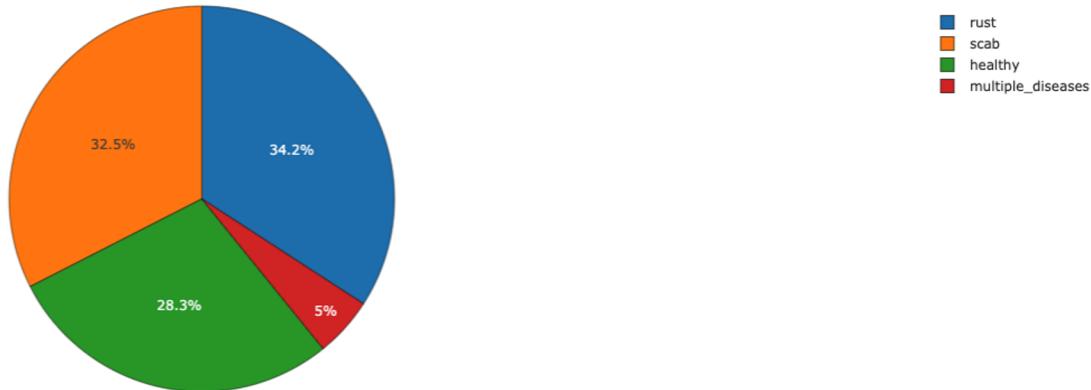
輸入資料分佈

Distribution of red channel values



輸入資料類別比例

Pie chart of targets



類神經模型結構 - InceptionResNetV2

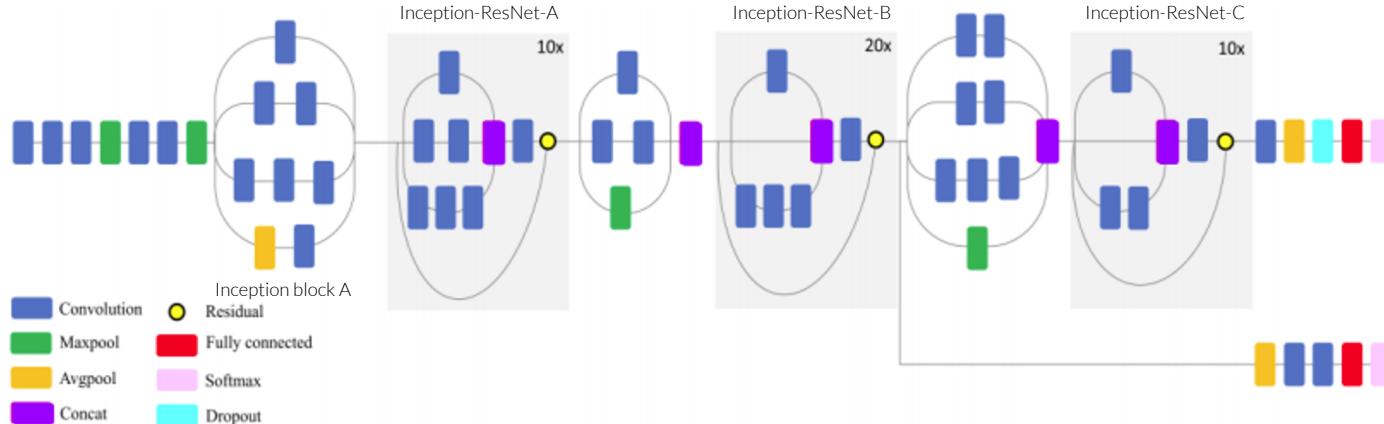
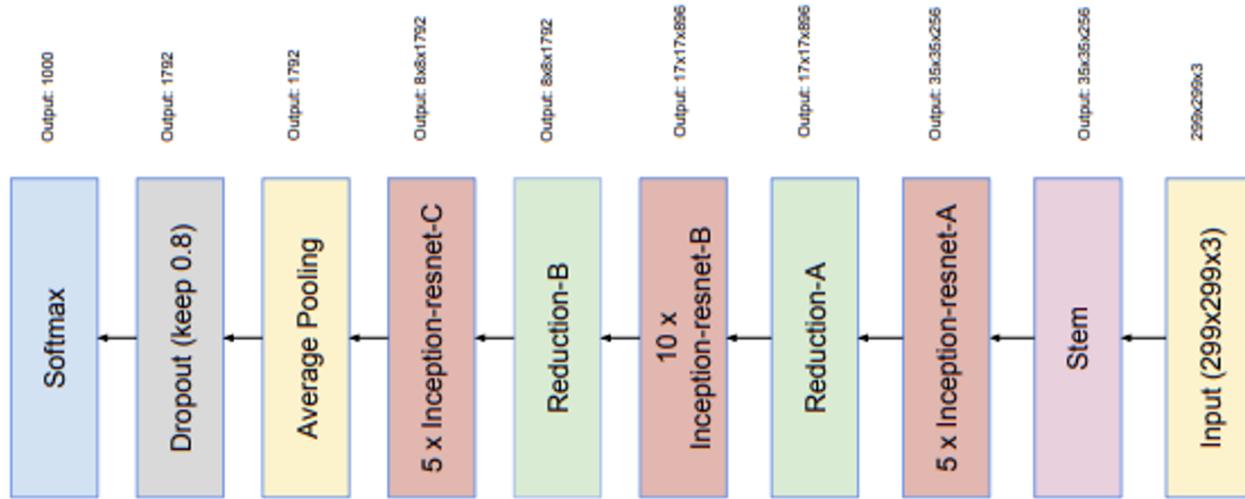


Figure 5. Schematic diagram of InceptionResNetV2 model (compressed view).

類神經模型結構 - InceptionResNetV2



類神經模型結構 - InceptionResNetV2

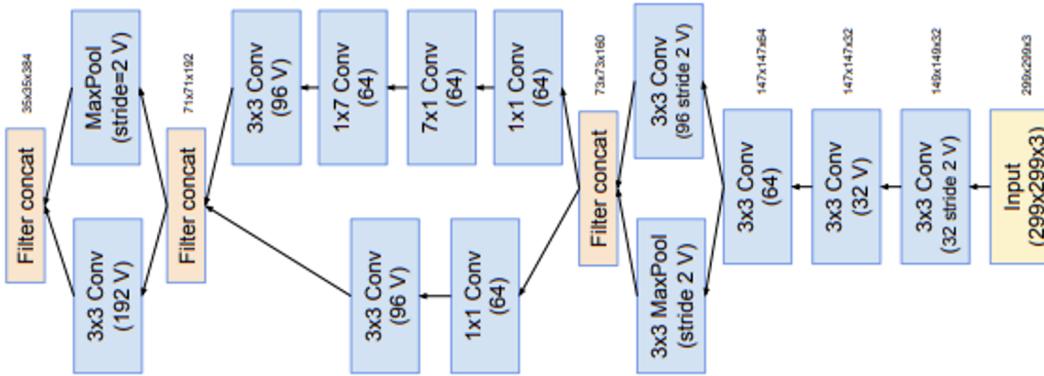
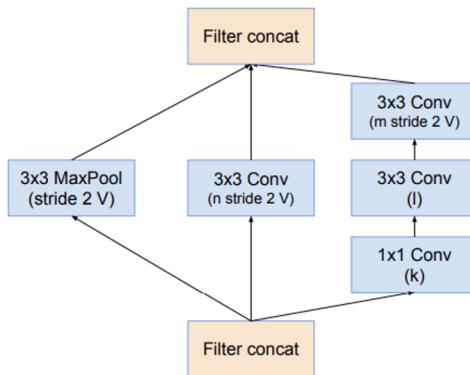


Figure 3. The schema for stem of the pure Inception-v4 and Inception-ResNet-v2 networks. This is the input part of those networks. Cf. Figures 9 and 15

類神經模型結構 - InceptionResNetV2

Inception block A



Inception block B

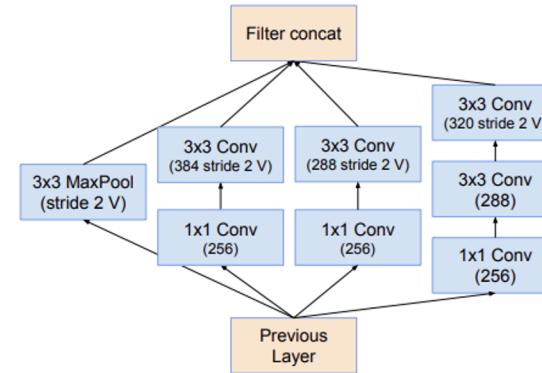


Figure 7. The schema for 35×35 to 17×17 reduction module. Different variants of this blocks (with various number of filters) are used in Figure 9, and 15 in each of the new Inception(-v4, -ResNet-v1, -ResNet-v2) variants presented in this paper. The k, l, m, n numbers represent filter bank sizes which can be looked up in Table 1.

Figure 18. The schema for 17×17 to 8×8 grid-reduction module. Reduction-B module used by the wider Inception-ResNet-v1 network in Figure 15.

類神經模型結構 - InceptionResNetV2

Inception-ResNet-A

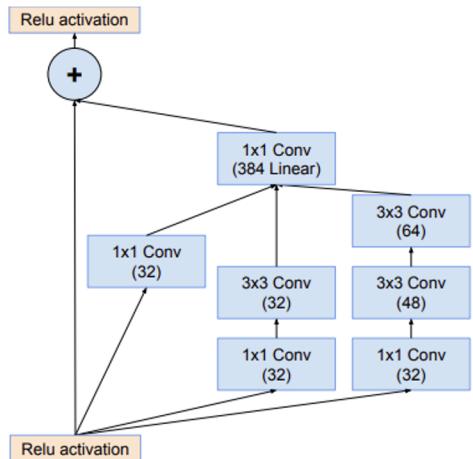


Figure 16. The schema for 35×35 grid (Inception-ResNet-A) module of the Inception-ResNet-v2 network.

Inception-ResNet-B

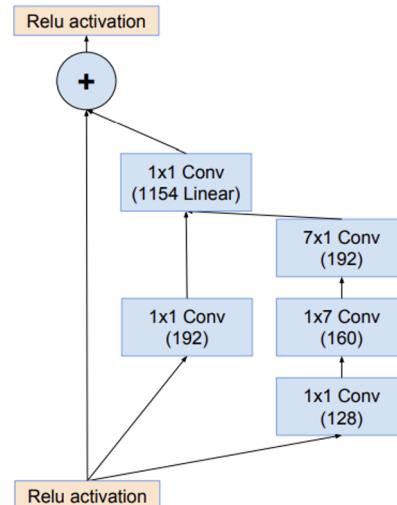


Figure 17. The schema for 17×17 grid (Inception-ResNet-B) module of the Inception-ResNet-v2 network.

Inception-ResNet-C

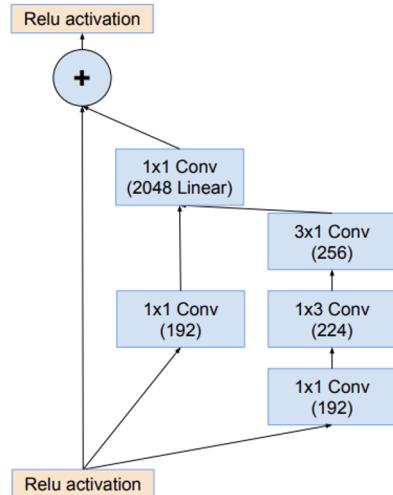


Figure 19. The schema for 8×8 grid (Inception-ResNet-C) module of the Inception-ResNet-v2 network.



相關參數 - InceptionResNetV2

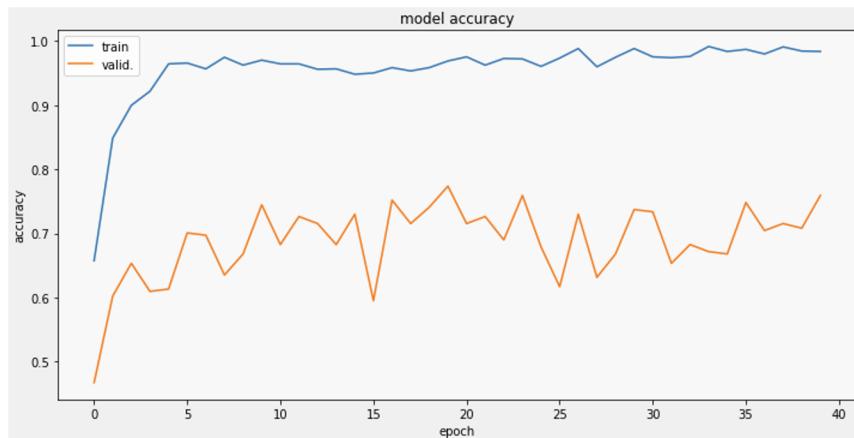
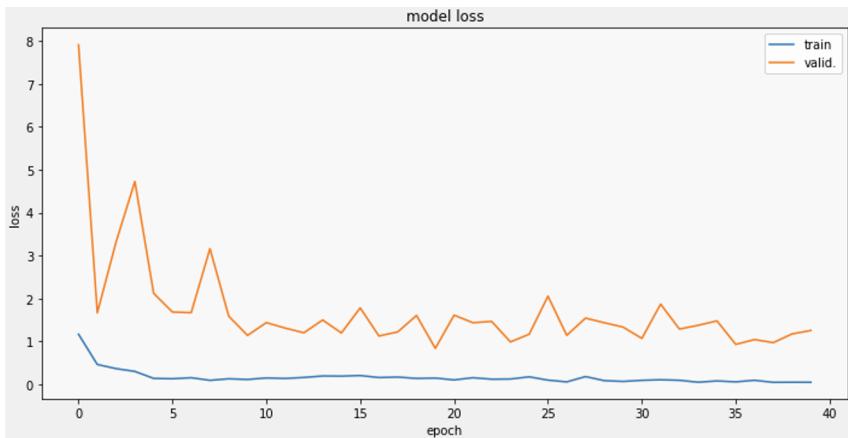
Layer (type)	Output Shape	Param #
inception_resnet_v2 (Model)	(None, 23, 23, 1536)	54336736
global_max_pooling2d (Global)	(None, 1536)	0
dense (Dense)	(None, 4)	6148
<hr/>		
Total params:	54,342,884	
Trainable params:	54,282,340	
Non-trainable params:	60,544	

output: 4個類別

Param: $6148 = (1536+1)*4$



結果 - InceptionResNetV2



類神經模型結構 - DenseNet

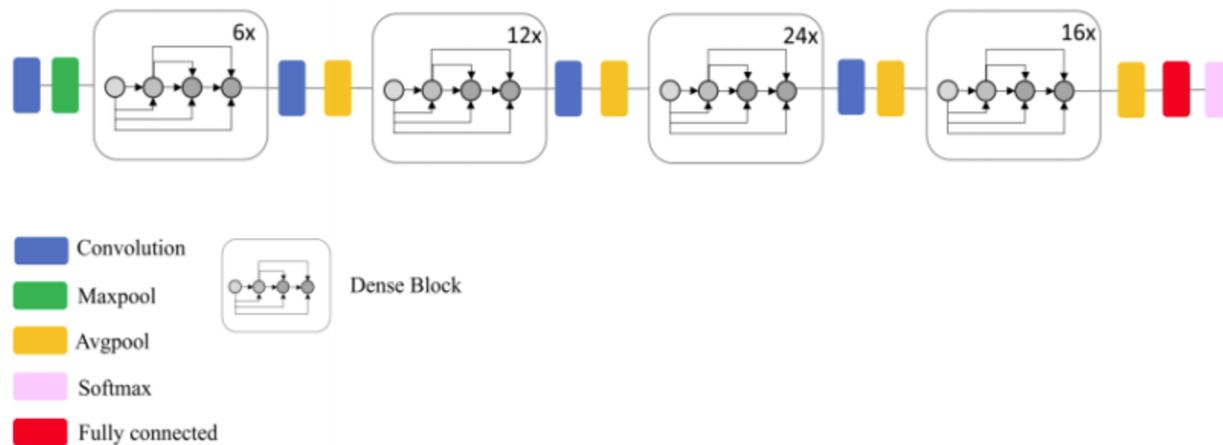


Figure 6. Schematic diagram of DenseNet model (compressed view).

類神經模型結構 - DenseNet

[4]

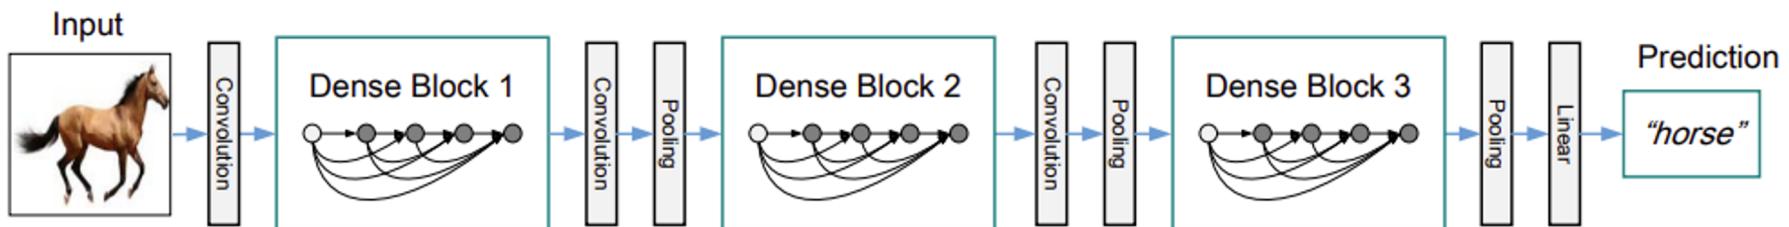


Figure 2: A deep DenseNet with three dense blocks. The layers between two adjacent blocks are referred to as transition layers and change feature-map sizes via convolution and pooling.

類神經模型結構 - DenseNet

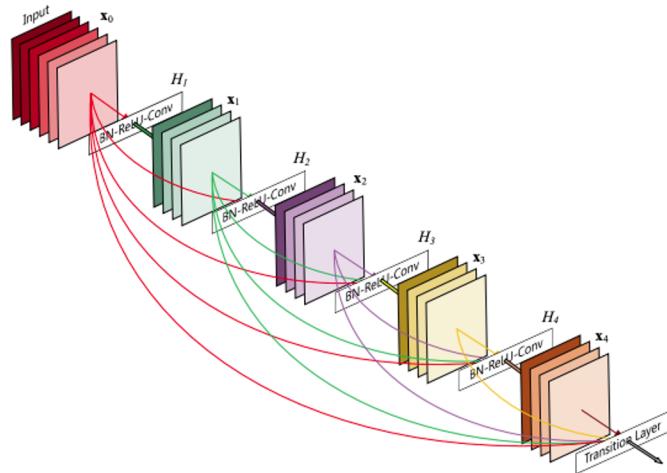


Figure 1: A 5-layer dense block with a growth rate of $k = 4$.
Each layer takes all preceding feature-maps as input.

我們用這

[4]

Layers	Output Size	DenseNet-121	DenseNet-169	DenseNet-201	DenseNet-264
Convolution	112 × 112				
Pooling	56 × 56				
Dense Block (1)	56 × 56	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$
Transition Layer (1)	56 × 56			$1 \times 1 \text{ conv}$	
	28 × 28			$2 \times 2 \text{ average pool, stride 2}$	
Dense Block (2)	28 × 28	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$
Transition Layer (2)	28 × 28			$1 \times 1 \text{ conv}$	
	14 × 14			$2 \times 2 \text{ average pool, stride 2}$	
Dense Block (3)	14 × 14	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 24$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 48$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 64$
Transition Layer (3)	14 × 14			$1 \times 1 \text{ conv}$	
	7 × 7			$2 \times 2 \text{ average pool, stride 2}$	
Dense Block (4)	7 × 7	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 16$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 48$
Classification Layer	1 × 1			$7 \times 7 \text{ global average pool}$	
				1000D fully-connected, softmax	



相關參數 - DenseNet

Layer (type)	Output Shape	Param #
densenet121 (Model)	(None, 16, 16, 1024)	7037504
global_average_pooling2d (G1	(None, 1024)	0
dense (Dense)	(None, 4)	4100
Total params:	7,041,604	
Trainable params:	6,957,956	
Non-trainable params:	83,648	

output: 4個類別

Param: $4100 = (1024+1)^*4$

相關參數 - DenseNet - modified

Layer (type)	Output Shape	Param #
densenet121 (Model)	(None, 16, 16, 1024)	7037504
global_average_pooling2d (G1	(None, 1024)	0
dense (Dense)	(None, 1024)	1049600
dense_1 (Dense)	(None, 1024)	1049600
dense_2 (Dense)	(None, 4)	4100

Total params: 9,140,804
Trainable params: 9,057,156
Non-trainable params: 83,648

output: 4個類別

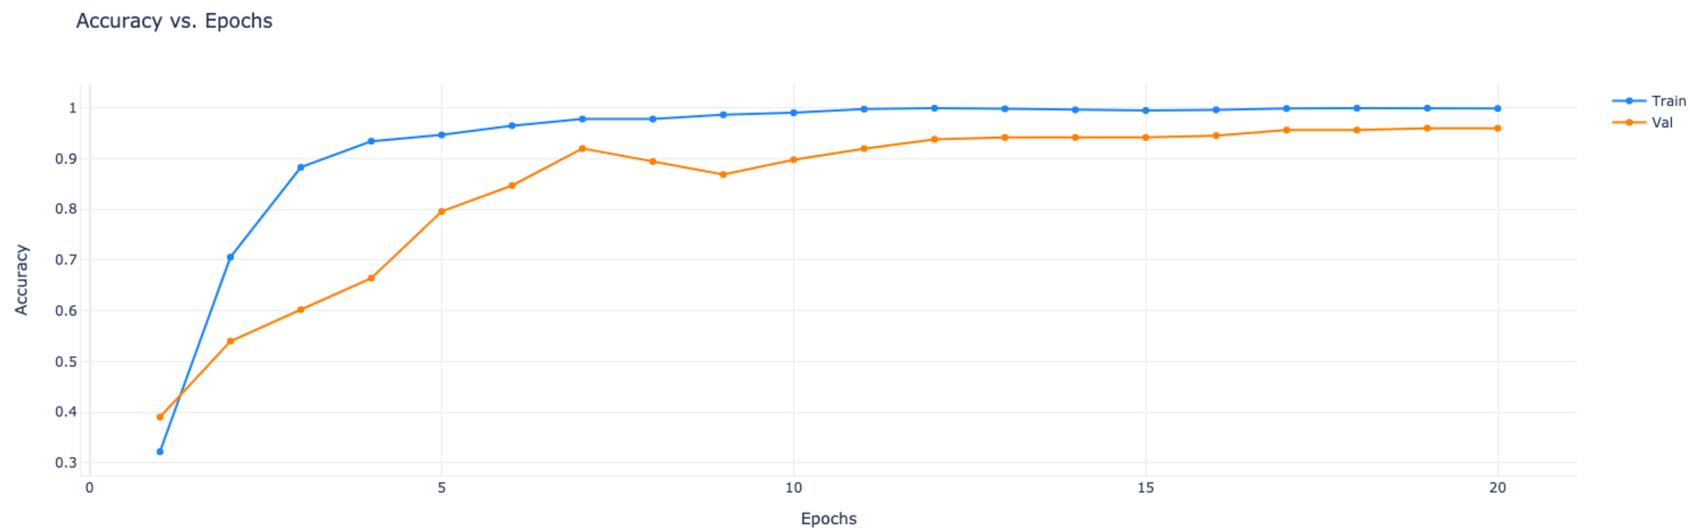
Param:

$$1049600 = (1024+1) * 1024$$

$$4100 = (1024+1) * 4$$

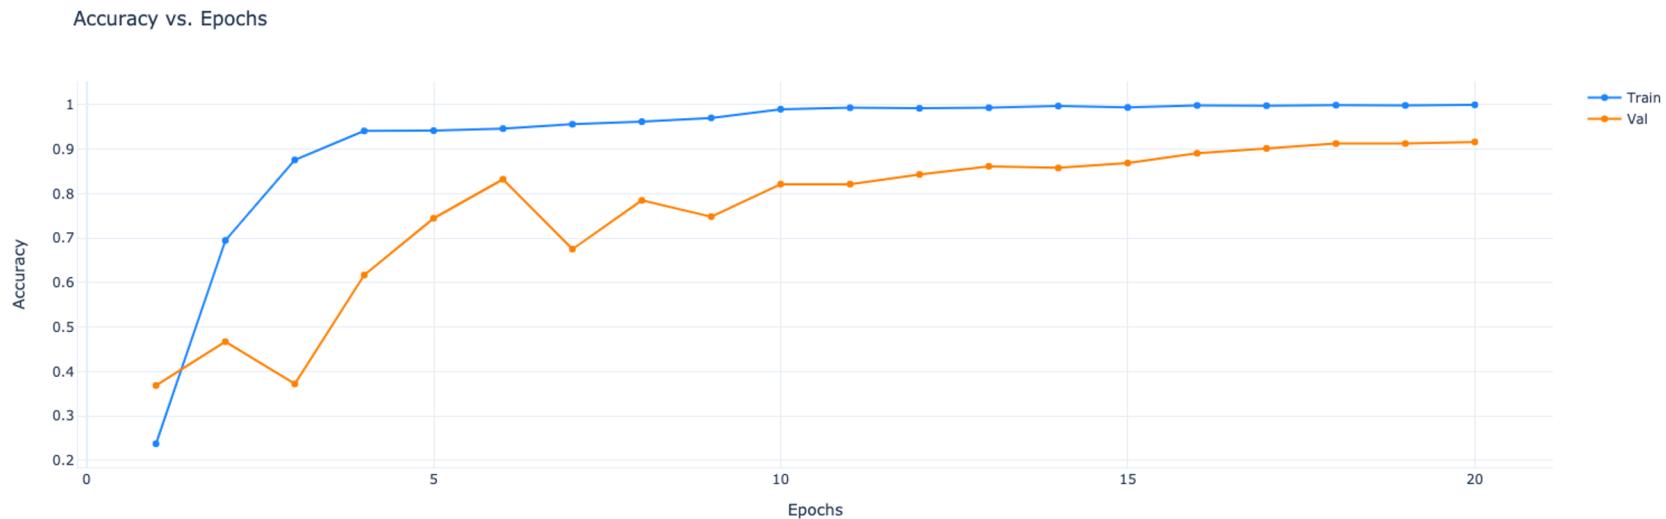


結果 - DenseNet

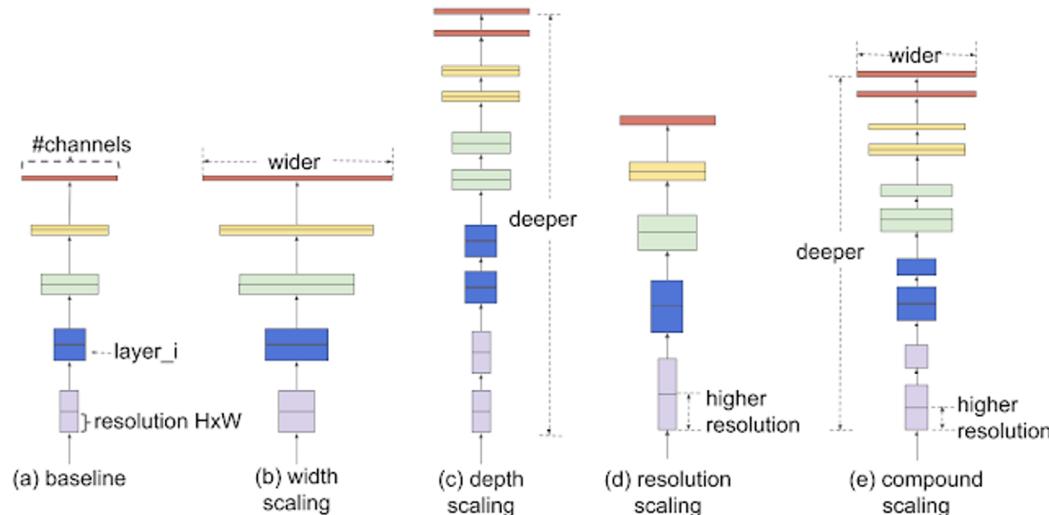




結果 - DenseNet - modified



類神經模型結構 - EfficientNet

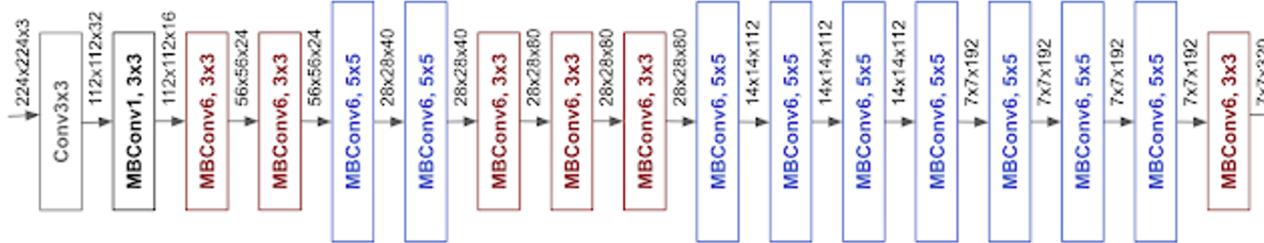


Comparison of different scaling methods. Unlike conventional scaling methods (b)-(d) that arbitrary scale a single dimension of the network, our compound scaling method uniformly scales up all dimensions in a principled way.



類神經模型結構 - EfficientNet-Bo

[5]



The architecture for our baseline network EfficientNet-B0 is simple and clean, making it easier to scale and generalize.

類神經模型結構 - EfficientNet-Bo^[5]

Table 1. EfficientNet-B0 baseline network – Each row describes a stage i with \hat{L}_i layers, with input resolution $\langle \hat{H}_i, \hat{W}_i \rangle$ and output channels \hat{C}_i . Notations are adopted from equation 2.

Stage i	Operator $\hat{\mathcal{F}}_i$	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels \hat{C}_i	#Layers \hat{L}_i
1	Conv3x3	224×224	32	1
2	MBConv1, k3x3	112×112	16	1
3	MBConv6, k3x3	112×112	24	2
4	MBConv6, k5x5	56×56	40	2
5	MBConv6, k3x3	28×28	80	3
6	MBConv6, k5x5	14×14	112	3
7	MBConv6, k5x5	14×14	192	4
8	MBConv6, k3x3	7×7	320	1
9	Conv1x1 & Pooling & FC	7×7	1280	1

$$\mathcal{N} = \bigodot_{i=1 \dots s} \mathcal{F}_i^{L_i}(X_{\langle H_i, W_i, C_i \rangle}) \quad (1)$$

where $\mathcal{F}_i^{L_i}$ denotes layer F_i is repeated L_i times in stage i , H_i, W_i, C_i denotes the shape of input tensor X of layer

$$\begin{aligned} \max_{d, w, r} \quad & \text{Accuracy}(\mathcal{N}(d, w, r)) \\ \text{s.t.} \quad & \mathcal{N}(d, w, r) = \bigodot_{i=1 \dots s} \hat{\mathcal{F}}_i^{d \cdot \hat{L}_i}(X_{\langle r \cdot \hat{H}_i, r \cdot \hat{W}_i, w \cdot \hat{C}_i \rangle}) \end{aligned} \quad (2)$$

$\text{Memory}(\mathcal{N}) \leq \text{target_memory}$

$\text{FLOPS}(\mathcal{N}) \leq \text{target_flops}$

where w, d, r are coefficients for scaling network width, depth, and resolution; $\hat{\mathcal{F}}_i, \hat{L}_i, \hat{H}_i, \hat{W}_i, \hat{C}_i$ are predefined parameters in baseline network (see Table 1 as an example).

depth: $d = \alpha^\phi$

width: $w = \beta^\phi$

resolution: $r = \gamma^\phi$

s.t. $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$

$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$

(3)

類神經模型結構 - EfficientNet

[5]

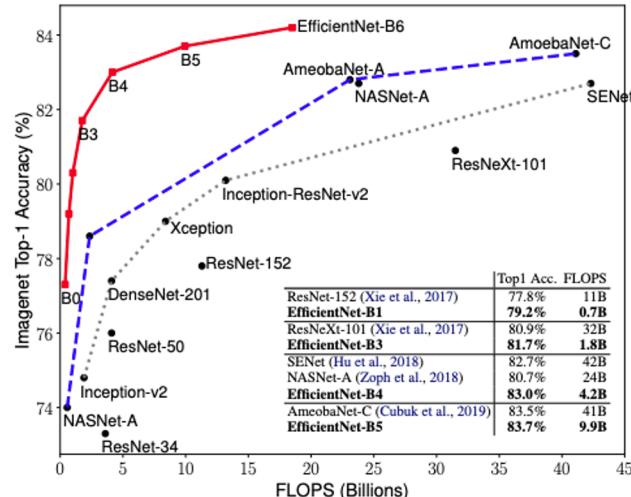


Figure 5. FLOPS vs. ImageNet Accuracy – Similar to Figure 1 except it compares FLOPS rather than model size.

相關參數 - EfficientNet

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
efficientnet-b7 (Model)	(None, 16, 16, 2560)	64097680
global_average_pooling2d_1 (Global Average Pooling)	(None, 2560)	0
dense_1 (Dense)	(None, 4)	10244
<hr/>		
Total params:	64,107,924	
Trainable params:	63,797,204	
Non-trainable params:	310,720	

output: 4個類別

Param: $10244 = (2560+1)*4$



相關參數 - EfficientNet-modified

Layer (type)	Output Shape	Param #
efficientnet-b7 (Model)	(None, 16, 16, 2560)	64097680
global_average_pooling2d_1 ((None, 2560)	0
dense_3 (Dense)	(None, 2560)	6556160
dense_4 (Dense)	(None, 2560)	6556160
dense_5 (Dense)	(None, 4)	10244
Total params:	77,220,244	
Trainable params:	76,909,524	
Non-trainable params:	310,720	

output: 4個類別

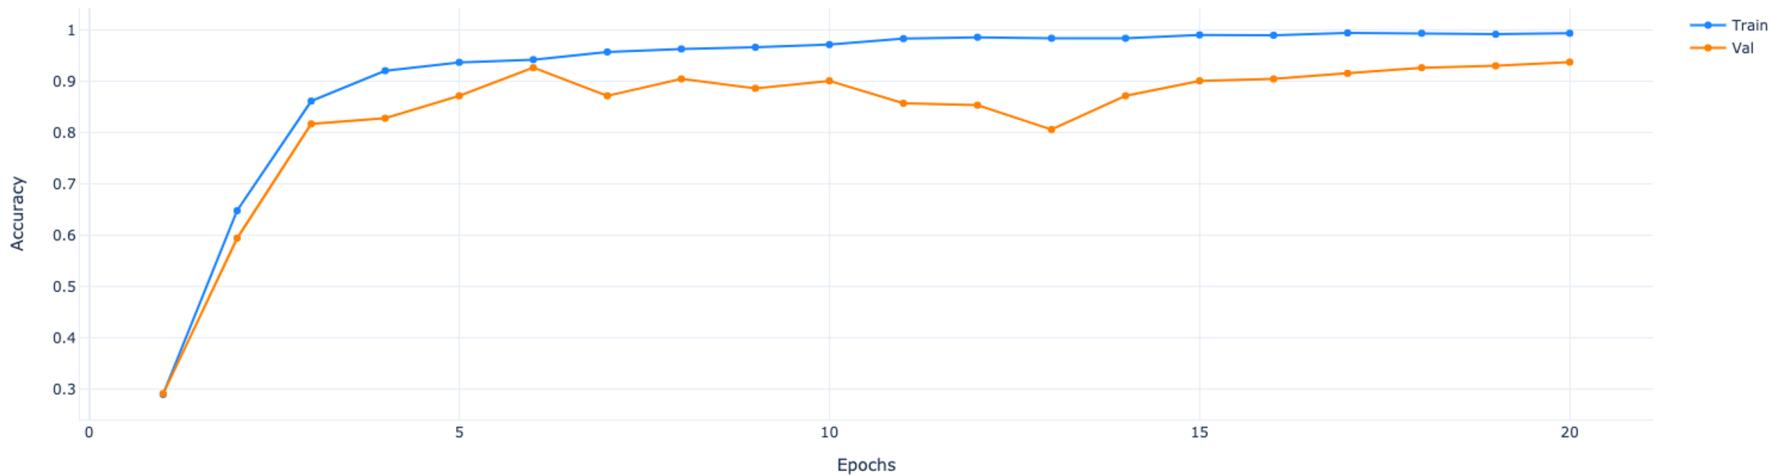
Param:

$$6556160 = (2560+1)*2560$$

$$10244 = (2560+1)*4$$

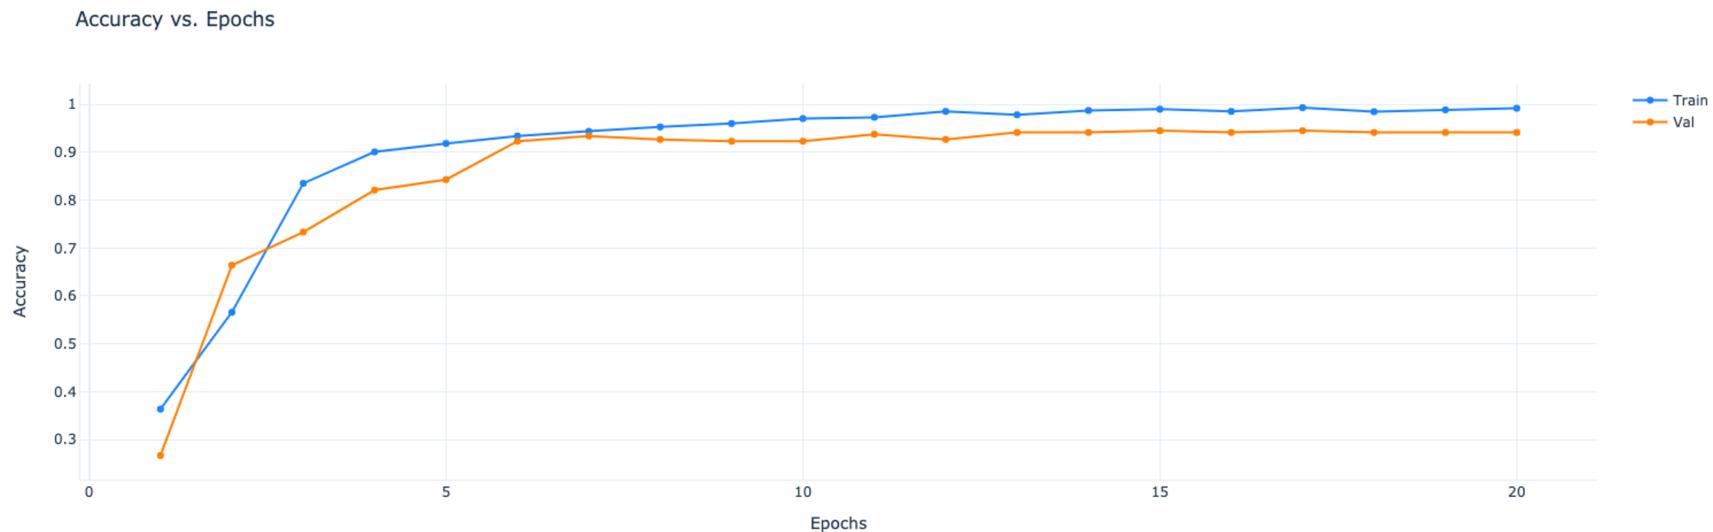


結果 - EfficientNet

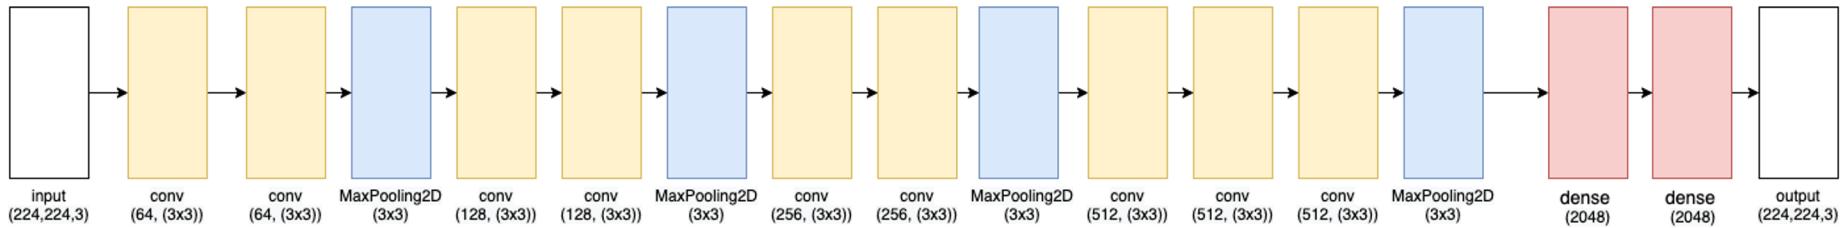




結果 - EfficientNet-modified



類神經模型結構 - leafNet (our model)



類神經模型結構 - leafNet與VGG19不同的地方



相關參數 - leafNet

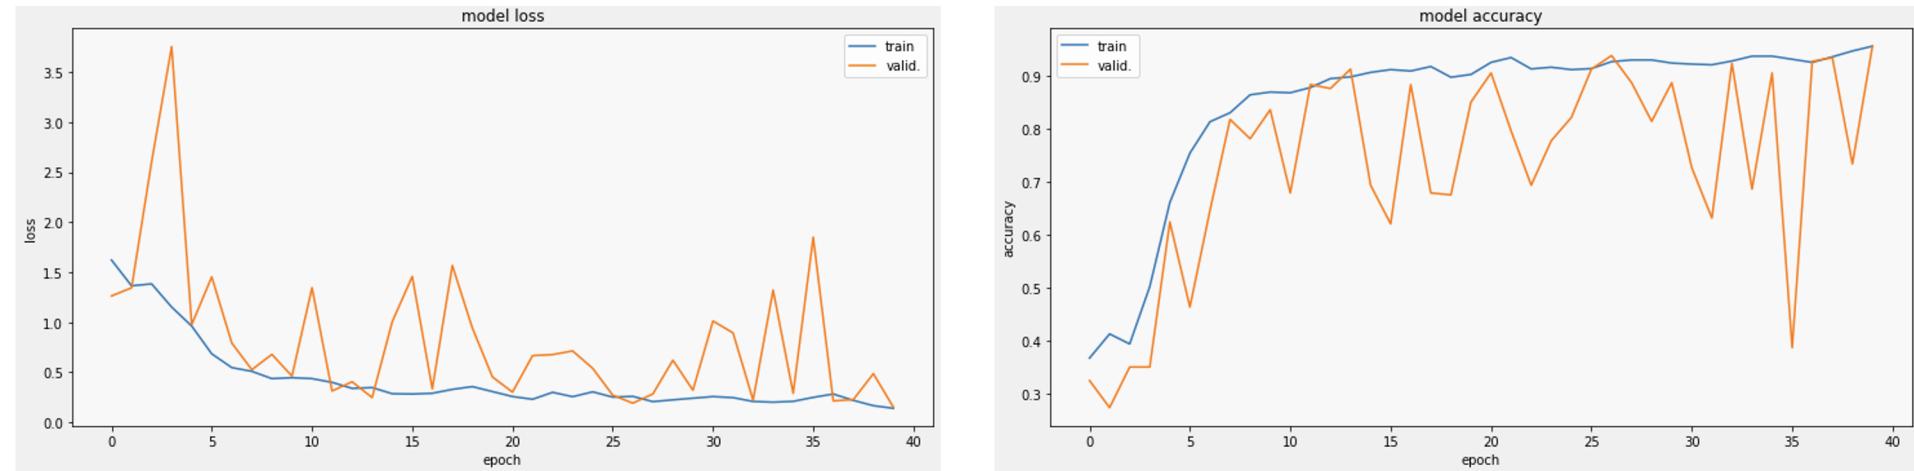
Model: "sequential_1"		
Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 224, 224, 64)	1792
batch_normalization_11 (Batch Normalization)	(None, 224, 224, 64)	256
activation_12 (Activation)	(None, 224, 224, 64)	0
conv2d_10 (Conv2D)	(None, 224, 224, 64)	36928
batch_normalization_12 (Batch Normalization)	(None, 224, 224, 64)	256
activation_13 (Activation)	(None, 224, 224, 64)	0
max_pooling2d_4 (MaxPooling2D)	(None, 74, 74, 64)	0
conv2d_11 (Conv2D)	(None, 74, 74, 128)	73856
batch_normalization_13 (Batch Normalization)	(None, 74, 74, 128)	512
activation_14 (Activation)	(None, 74, 74, 128)	0
conv2d_12 (Conv2D)	(None, 74, 74, 128)	147584
batch_normalization_14 (Batch Normalization)	(None, 74, 74, 128)	512
activation_15 (Activation)	(None, 74, 74, 128)	0
max_pooling2d_5 (MaxPooling2D)	(None, 24, 24, 128)	0
conv2d_13 (Conv2D)	(None, 24, 24, 256)	295168

conv2d_13 (Conv2D)	(None, 24, 24, 256)	295168
batch_normalization_15 (Batch Normalization)	(None, 24, 24, 256)	1024
activation_16 (Activation)	(None, 24, 24, 256)	0
conv2d_14 (Conv2D)	(None, 24, 24, 256)	590080
batch_normalization_16 (Batch Normalization)	(None, 24, 24, 256)	1024
activation_17 (Activation)	(None, 24, 24, 256)	0
max_pooling2d_6 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_15 (Conv2D)	(None, 8, 8, 512)	1180160
batch_normalization_17 (Batch Normalization)	(None, 8, 8, 512)	2048
activation_18 (Activation)	(None, 8, 8, 512)	0
conv2d_16 (Conv2D)	(None, 8, 8, 512)	2359808
batch_normalization_18 (Batch Normalization)	(None, 8, 8, 512)	2048
activation_19 (Activation)	(None, 8, 8, 512)	0
conv2d_17 (Conv2D)	(None, 8, 8, 512)	2359808
batch_normalization_19 (Batch Normalization)	(None, 8, 8, 512)	2048
activation_20 (Activation)	(None, 8, 8, 512)	0
max_pooling2d_7 (MaxPooling2D)	(None, 2, 2, 512)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_3 (Dense)	(None, 2048)	4196352

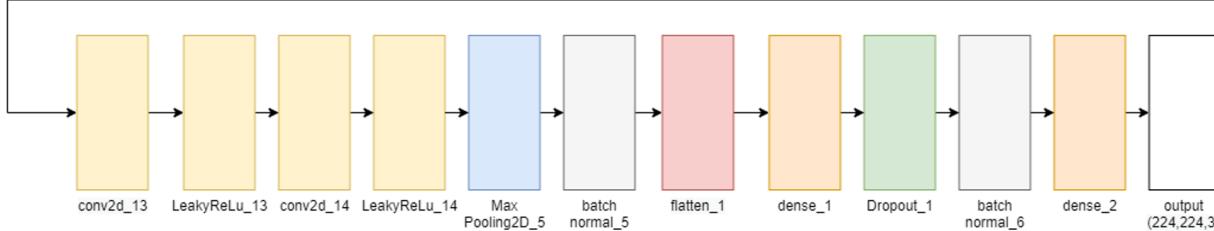
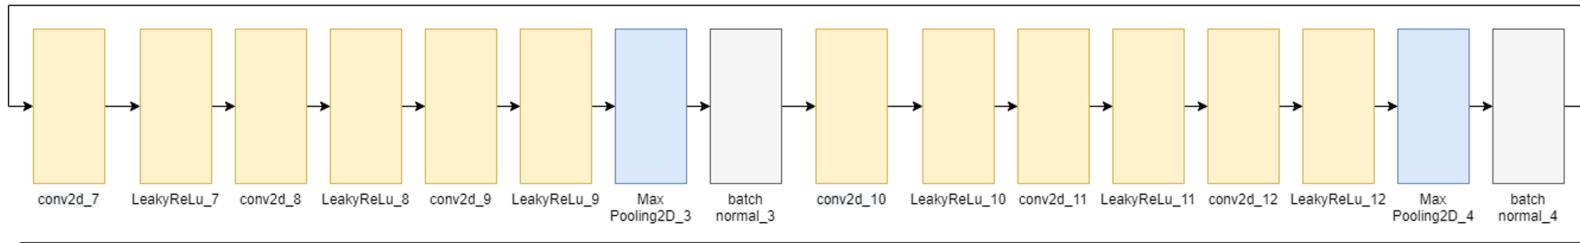
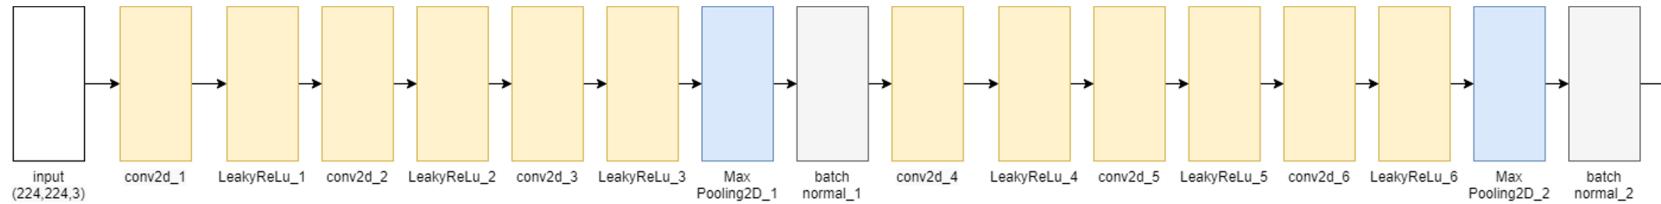
dense_3 (Dense)	(None, 2048)	4196352
batch_normalization_20 (Batch Normalization)	(None, 2048)	8192
activation_21 (Activation)	(None, 2048)	0
dense_4 (Dense)	(None, 2048)	4196352
batch_normalization_21 (Batch Normalization)	(None, 2048)	8192
activation_22 (Activation)	(None, 2048)	0
dense_5 (Dense)	(None, 4)	8196
activation_23 (Activation)	(None, 4)	0
<hr/>		
Total params:	15,472,196	
Trainable params:	15,459,140	
Non-trainable params:	13,056	



結果 - leafNet



相關參數 - model - 2



相關參數 - model - 2

Using TensorFlow backend.
Model: "sequential_1"

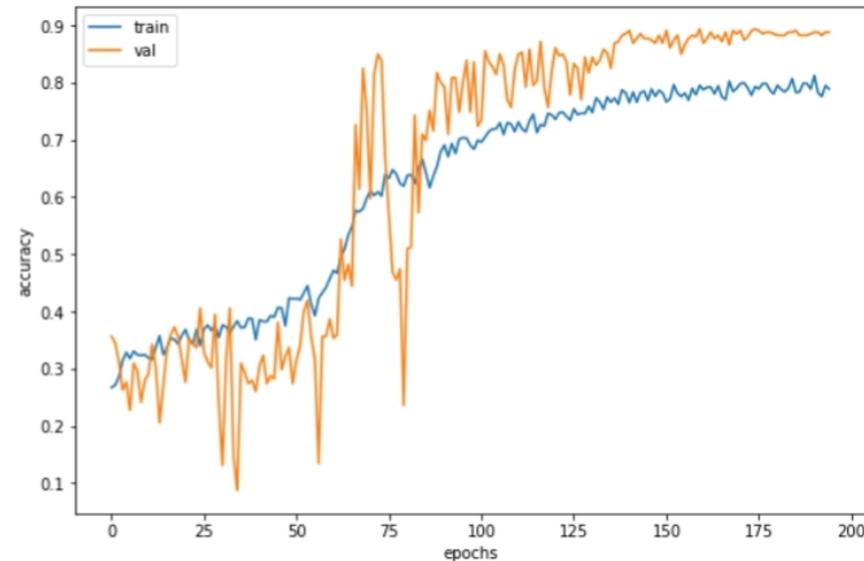
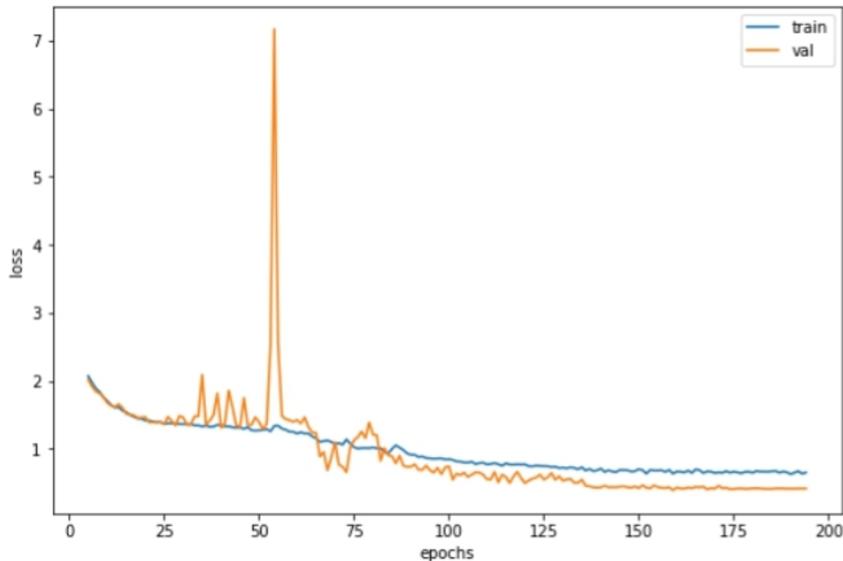
Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 222, 222, 32)	896
leaky_re_lu_1 (LeakyReLU)	(None, 222, 222, 32)	0
conv2d_2 (Conv2D)	(None, 220, 220, 32)	9248
leaky_re_lu_2 (LeakyReLU)	(None, 220, 220, 32)	0
conv2d_3 (Conv2D)	(None, 216, 216, 32)	25632
leaky_re_lu_3 (LeakyReLU)	(None, 216, 216, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 108, 108, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 108, 108, 32)	128
conv2d_4 (Conv2D)	(None, 106, 106, 64)	18496
leaky_re_lu_4 (LeakyReLU)	(None, 106, 106, 64)	0
conv2d_5 (Conv2D)	(None, 104, 104, 64)	36928
leaky_re_lu_5 (LeakyReLU)	(None, 104, 104, 64)	0
conv2d_6 (Conv2D)	(None, 100, 100, 64)	102464

[] conv2d_6 (Conv2D)	(None, 100, 100, 64)	102464
↳ leaky_re_lu_6 (LeakyReLU)	(None, 100, 100, 64)	0
max_pooling2d_2 (MaxPooling2D)	(None, 50, 50, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 50, 50, 64)	256
conv2d_7 (Conv2D)	(None, 48, 48, 128)	73856
leaky_re_lu_7 (LeakyReLU)	(None, 48, 48, 128)	0
conv2d_8 (Conv2D)	(None, 46, 46, 128)	147584
leaky_re_lu_8 (LeakyReLU)	(None, 46, 46, 128)	0
conv2d_9 (Conv2D)	(None, 42, 42, 128)	409728
leaky_re_lu_9 (LeakyReLU)	(None, 42, 42, 128)	0
max_pooling2d_3 (MaxPooling2D)	(None, 21, 21, 128)	0
batch_normalization_3 (Batch Normalization)	(None, 21, 21, 128)	512
conv2d_10 (Conv2D)	(None, 19, 19, 256)	295168
leaky_re_lu_10 (LeakyReLU)	(None, 19, 19, 256)	0
conv2d_11 (Conv2D)	(None, 17, 17, 256)	590080
leaky_re_lu_11 (LeakyReLU)	(None, 17, 17, 256)	0
conv2d_12 (Conv2D)	(None, 13, 13, 256)	1638656
leaky_re_lu_12 (LeakyReLU)	(None, 13, 13, 256)	0
max_pooling2d_4 (MaxPooling2D)	(None, 6, 6, 256)	0

max_pooling2d_4 (MaxPooling2D)	(None, 6, 6, 256)	0
batch_normalization_4 (Batch Normalization)	(None, 6, 6, 256)	1024
conv2d_13 (Conv2D)	(None, 4, 4, 512)	1180160
leaky_re_lu_13 (LeakyReLU)	(None, 4, 4, 512)	0
conv2d_14 (Conv2D)	(None, 2, 2, 512)	2359808
leaky_re_lu_14 (LeakyReLU)	(None, 2, 2, 512)	0
max_pooling2d_5 (MaxPooling2D)	(None, 1, 1, 512)	0
batch_normalization_5 (Batch Normalization)	(None, 1, 1, 512)	2048
flatten_1 (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 16)	8208
dropout_1 (Dropout)	(None, 16)	0
batch_normalization_6 (Batch Normalization)	(None, 16)	64
dense_2 (Dense)	(None, 4)	68
<hr/>		
Total params:	6,901,012	
Trainable params:	6,898,996	
Non-trainable params:	2,016	



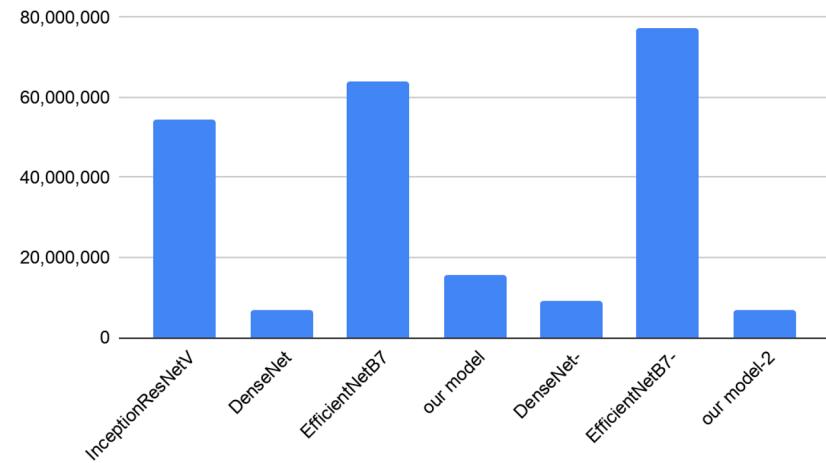
結果 - model - 2



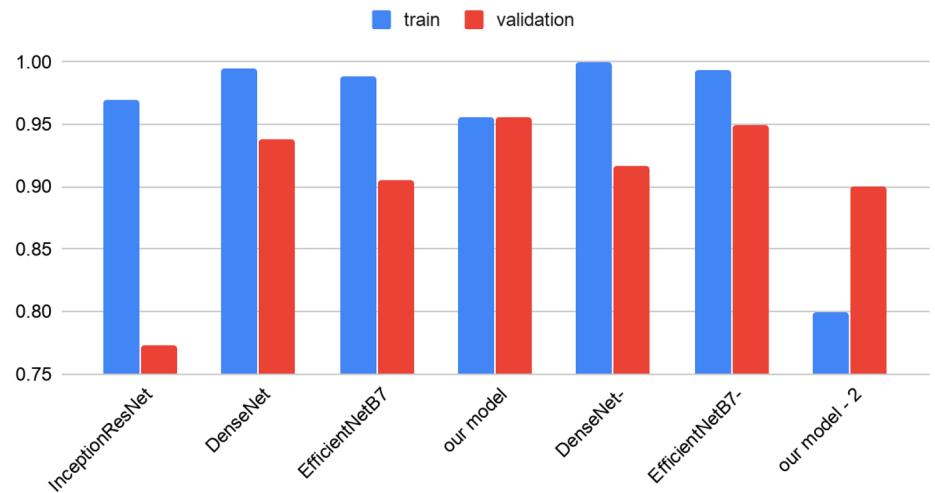


綜合比較

縱軸: params, 橫軸:

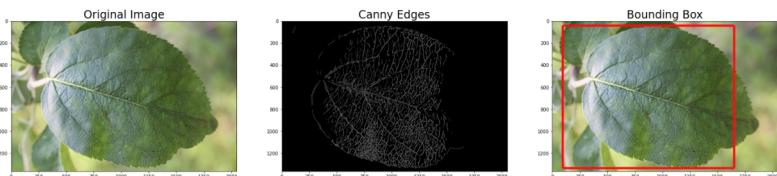
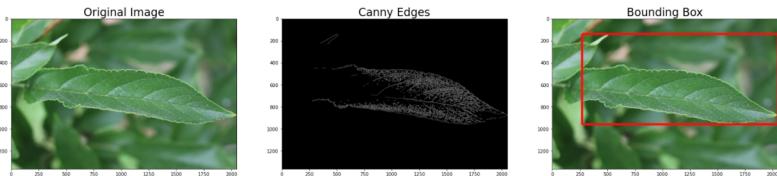
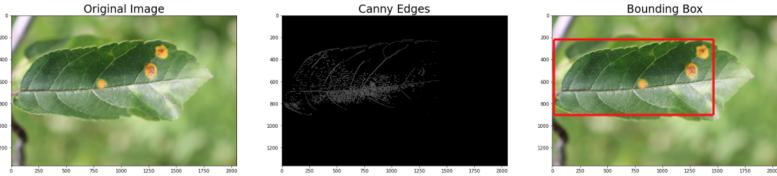


train和validation



[7]

可以改進的地方



(a)



(b)

Fig. 12: (a) Image inputted, (b) Image result

實作講解



Distributed training with TensorFlow

定義 strategy

```
try:  
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver()  
    print('Running on TPU ', tpu.master())  
except ValueError:  
    tpu = None  
  
if tpu:  
    tf.config.experimental_connect_to_cluster(tpu)  
    tf.tpu.experimental.initialize_tpu_system(tpu)  
    strategy = tf.distribute.experimental.TPUStrategy(tpu)  
    print('use tpu')  
else:  
    if tf.config.list_physical_devices('GPU'):  
        strategy = tf.distribute.MirroredStrategy()  
        print('use gpu')  
    else: # use default strategy  
        print('use cpu')  
    strategy = tf.distribute.get_strategy()
```

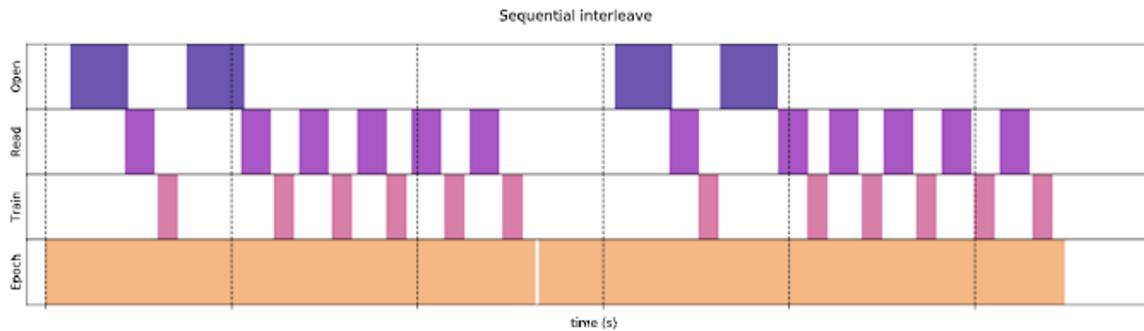
使用

```
In [16]: with strategy.scope():  
    model3 = Sequential()  
  
    model3.add(Conv2D(64, (3, 3), input_shape=(28, 28, 1)))  
    model3.add(BatchNormalization(axis=-1))  
    model3.add(Activation('relu'))  
    model3.add(Conv2D(64, (3, 3), padding='same'))  
    model3.add(BatchNormalization(axis=-1))  
    model3.add(Activation('relu'))  
    model3.add(MaxPooling2D(pool_size=(3, 3)))
```

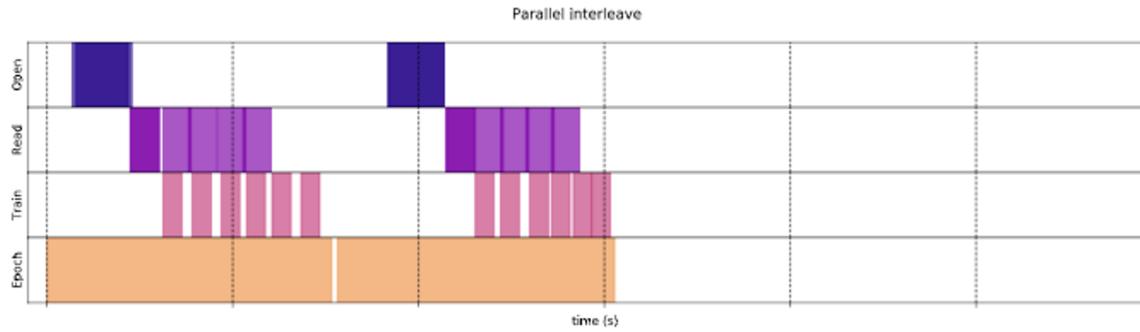
小提示 : tf.config.list_physical_devices()

Build TensorFlow input pipelines

origin



parallelized



Build TensorFlow input pipelines

定義 tf.data.Dataset

```
train_dataset_1 = (
    tf.data.Dataset
        .from_tensor_slices((train_paths, train_labels))
        .map(decode_image, num_parallel_calls=AUTO)
        .cache()
        .map(data_augment, num_parallel_calls=AUTO)
        .repeat()
        .shuffle(512)
        .batch(BATCH_SIZE)
        .prefetch(AUTO)
)
valid_dataset = (
    tf.data.Dataset
        .from_tensor_slices((valid_paths, valid_labels))
        .map(decode_image, num_parallel_calls=AUTO)
        .batch(BATCH_SIZE)
        .cache()
        .prefetch(AUTO)
)
test_dataset = (
    tf.data.Dataset
        .from_tensor_slices(test_paths)
        .map(decode_image, num_parallel_calls=AUTO)
        .map(data_augment, num_parallel_calls=AUTO)
        .batch(BATCH_SIZE)
)
```

使用

```
history3 = model3.fit(
    train_dataset_1,
    epochs=EPOCHS,
    callbacks=[es],
    steps_per_epoch=STEPS PER EPOCH,
    validation_data=valid_dataset if SPLIT_VALIDATION else None,
)
```

save model / load model

save

```
!mkdir -p saved_model  
model3.save('saved_model/our-model')  
import pandas as pd  
pd.DataFrame.from_dict(history3.history).to_csv('history-1.csv', index=False)  
  
INFO:tensorflow:Assets written to: saved_model/our-model/assets
```

小提示: 查看model存不存在

```
!ls saved_model  
!ls saved_model/model-1
```

```
model-1  model-1-1  
assets  saved_model.pb  variables
```

load

```
: InceptionResNetV2 = tf.keras.models.load_model('saved_model/model-1')  
InceptionResNetV2.summary()  
  
Current values:  
NotebookApp.iopub_msg_rate_limit=1000.0 (msgs/sec)  
NotebookApp.rate_limit_window=3.0 (secs)  
  
Model: "sequential"  


| Layer (type)                  | Output Shape         | Param #  |
|-------------------------------|----------------------|----------|
| inception_resnet_v2 (Model)   | (None, 23, 23, 1536) | 54336736 |
| global_max_pooling2d (Global) | (None, 1536)         | 0        |
| dense (Dense)                 | (None, 4)            | 6148     |
| Total params:                 | 54,342,884           |          |
| Trainable params:             | 54,282,340           |          |
| Non-trainable params:         | 60,544               |          |


```



相關連結

1. <https://www.kaggle.com/alanhc/fork-of-plant-pathology-2020-eda-models?scriptVersionId=37593370>
2. <https://www.kaggle.com/alanhc/fork-of-plant-pathology-2020-eda-models?scriptVersionId=37674876>
3. <http://120.125.83.237:8086/notebooks/demo.ipynb>
4. <http://120.125.83.237:8086/notebooks/week18%20-%20IncepresnetV2%2BENB7-Copy1.ipynb>
5. <http://120.125.83.237:8086/notebooks/week18%20-%20IncepresnetV2%2BENB7-Copy2.ipynb>
6. <https://colab.research.google.com/drive/1WBNOQE1CzBEuRrvgKjlPUACdK-eYPzFk?usp=sharing>
7. <https://docs.google.com/presentation/d/1S1qe0ZPdjJQNO04-AOHunQ4qWvtONIWou6CSfR1ENGs/edit?usp=sharing>
8. <https://docs.google.com/document/d/1G7mgdnRqBf98OWn2IThESGO2uYmfNhpiA8Fc0BKG4U0/edit?usp=sharing>

附錄 - efficientNet計算-以B0為例

Table 1. EfficientNet-B0 baseline network – Each row describes a stage i with \hat{L}_i layers, with input resolution $\langle \hat{H}_i, \hat{W}_i \rangle$ and output channels \hat{C}_i . Notations are adopted from equation 2.

Stage i	Operator $\hat{\mathcal{F}}_i$	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels \hat{C}_i	#Layers \hat{L}_i
1	Conv3x3	224×224	32	1
2	MBCConv1, k3x3	112×112	16	1
3	MBCConv6, k3x3	112×112	24	2
4	MBCConv6, k5x5	56×56	40	2
5	MBCConv6, k3x3	28×28	80	3
6	MBCConv6, k5x5	14×14	112	3
7	MBCConv6, k5x5	14×14	192	4
8	MBCConv6, k3x3	7×7	320	1
9	Conv1x1 & Pooling & FC	7×7	1280	1

Starting from the baseline EfficientNet-B0, we apply our compound scaling method to scale it up with two steps:

- STEP 1: we first fix $\phi = 1$, assuming twice more resources available, and do a small grid search of α, β, γ based on Equation 2 and 3. In particular, we find the best values for EfficientNet-B0 are $\alpha = 1.2, \beta = 1.1, \gamma = 1.15$, under constraint of $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$.
- STEP 2: we then fix α, β, γ as constants and scale up baseline network with different ϕ using Equation 3, to obtain EfficientNet-B1 to B7 (Details in Table 2).

Notably, it is possible to achieve even better performance by searching for α, β, γ directly around a large model, but the search cost becomes prohibitively more expensive on larger models. Our method solves this issue by only doing search once on the small baseline network (step 1), and then use the same scaling coefficients for all other models (step 2).



參考

1. Thapa, R., Snavely, N., Belongie, S., & Khan, A. (2020). The Plant Pathology 2020 challenge dataset to classify foliar disease of apples. *arXiv preprint arXiv:2004.11958*.
2. Mahdianpari, M., Salehi, B., Rezaee, M., Mohammadimanesh, F., & Zhang, Y. (2018). Very deep convolutional neural networks for complex land cover mapping using multispectral remote sensing imagery. *Remote Sensing*, 10(7), 1119.
3. Szegedy, C., Ioffe, S., Vanhoucke, V., & Alemi, A. A. (2017, February). Inception-v4, inception-resnet and the impact of residual connections on learning. In *Thirty-first AAAI conference on artificial intelligence*.
4. Huang, G., Liu, Z., Van Der Maaten, L., & Weinberger, K. Q. (2017). Densely connected convolutional networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 4700-4708).
5. Tan, M., & Le, Q. V. (2019). Efficientnet: Rethinking model scaling for convolutional neural networks. *arXiv preprint arXiv:1905.11946*.
6. Tan, M., Chen, B., Pang, R., Vasudevan, V., Sandler, M., Howard, A., & Le, Q. V. (2019). Mnasnet: Platform-aware neural architecture search for mobile. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (pp. 2820-2828).
7. Veites-Campos, S. A., Ramírez-Betancour, R., & González-Pérez, M. (2018). Identification of cocoa pods with image processing and artificial neural networks. *International Journal of Advanced Engineering, Management and Science*, 4(7).



參考

8. Mazen, F. M., & Nashat, A. A. (2019). Ripeness classification of bananas using an artificial neural network. *Arabian Journal for Science and Engineering*, 44(8), 6901-6910.
9. Hassan, N. M. H., & Nashat, A. A. (2019). New effective techniques for automatic detection and classification of external olive fruits defects based on image processing techniques. *Multidimensional Systems and Signal Processing*, 30(2), 571-589.
10. Basri, H., Syarif, I., Sukardhoto, S., & Falah, M. F. (2019). INTELLIGENT SYSTEM FOR AUTOMATIC CLASSIFICATION OF FRUIT DEFECT USING FASTER REGION-BASED CONVOLUTIONAL NEURAL NETWORK (FASTER R-CNN). *Jurnal Ilmiah Kursor*, 10(1).
11. Sahu, D., & Dewangan, C. (2017). Identification and classification of mango fruits using image processing. *Int. J. Sci. Res. Comput. Sci. Eng. Inf. Technol*, 2(2), 203-210.
12. Sahu, D., & Pottdar, R. M. (2017). Defect identification and maturity detection of mango fruits using image analysis. *American Journal of Artificial Intelligence*, 1(1), 5-14.
13. Thendral, R., & Suhasini, A. (2017). Automated skin defect identification system for orange fruit grading based on genetic algorithm. *Current Sci*, 112(8), 1704-1711.
14. Ireri, D., Belal, E., Okinda, C., Makange, N., & Ji, C. (2019). A computer vision system for defect discrimination and grading in tomatoes using machine learning and image processing. *Artificial Intelligence in Agriculture*, 2, 28-37.